

1                    Harnessing Large Language Models for Adaptive and  
2                    Explainable Traffic Forecasting

3                    Haoyang Yan<sup>1</sup> and Xiaolei Ma<sup>\*1,2</sup>

4                    <sup>1</sup>School of Transportation Science and Engineering, Beihang University, Beijing,  
5                    China.

6                    <sup>2</sup>Key Laboratory of Intelligent Transportation Technology and System of the  
7                    Ministry of Education, Beihang University, Beijing 102206, China

8                    \*Address correspondence to: xiaolei@buaa.edu.cn

9                    **Abstract**

10                  Accurate traffic prediction is fundamental for Intelligent Transportation Systems (ITS) aiming  
11                  to alleviate congestion and improve urban mobility resilience. While deep learning ap-  
12                  proaches like graph neural networks and sequence models have advanced short-term forecasting  
13                  accuracy under standard conditions, their practical deployment is often hampered by noticeable  
14                  limitations. These models typically struggle to generalize during anomalous events such as road  
15                  incidents or severe weather, exhibit inflexibility in incorporating diverse real-time contextual  
16                  information, and produce outputs that lack the interpretability and actionable insights crucial  
17                  for operational traffic management. Separately, large language models (LLMs), despite their  
18                  reasoning power, face inherent challenges in direct numerical time-series prediction and require  
19                  substantial resources for task-specific fine-tuning, limiting their standalone applicability. Here,  
20                  we introduce Chat-ITS, a novel hybrid framework designed to overcome these challenges by syn-  
21                  ergistically combining robust probabilistic time-series forecasting with the contextual reasoning  
22                  capabilities of LLMs. Chat-ITS first generates multiple candidate traffic state trajectories and  
23                  associated uncertainty bounds using a specialized probabilistic forecaster. Subsequently, an  
24                  LLM processes these candidates, conditioned on flexible natural language prompts that encode  
25                  both structured data (e.g., weather forecasts, road closures) and unstructured information (e.g.,  
26                  descriptions of public events). The LLM selects the most contextually plausible trajectory and,  
27                  critically, generates human-readable explanations and actionable recommendations for traffic  
28                  operators. We demonstrate through extensive evaluation under both routine and diverse anomalous  
29                  scenarios that Chat-ITS enhances prediction accuracy during irregular events compared to  
30                  baseline models, while maintaining state-of-the-art performance under normal traffic conditions.  
31                  Furthermore, case studies highlight the framework's ability to handle novel events described only  
32                  through text prompts by leveraging the LLM's reasoning to select the most plausible outcome  
33                  from a range of probabilistically generated forecasts; this provides context-aware, actionable in-  
34                  sights (e.g., suggesting specific signal timing adjustments or dynamic routing strategies), thereby

35 bridging the gap towards more adaptive, effective, and practical ITS applications.

## 36 1 Introduction

37 Accurate traffic prediction is fundamental to the efficacy of Intelligent Transportation Systems  
38 (ITS), enabling critical functions such as dynamic route guidance, adaptive traffic signal control,  
39 and proactive incident management essential for mitigating congestion, reducing emissions, and  
40 enhancing urban mobility resilience [1]. Congestion alone costs economies billions annually and  
41 degrades quality of life in urban centers [2]. Effective ITS, powered by reliable forecasts, promises  
42 substantial improvements in transportation efficiency and sustainability. Recent advances, partic-  
43 ularly the application of deep learning techniques like graph neural networks (GNNs) for modeling  
44 complex spatial dependencies across road networks [3] and sophisticated sequence models (e.g., tem-  
45 poral convolution networks, attention mechanisms) for capturing temporal dynamics [4, 5], have  
46 considerable improved short-term forecasting accuracy under typical, recurring traffic conditions [6].  
47 These methods effectively learn patterns from large historical datasets, providing a strong foundation  
48 for next-generation ITS applications operating under predictable circumstances.

49 Despite these successes, existing state-of-the-art traffic forecasting methods face critical limita-  
50 tions that hinder their real-world operational utility, particularly under non-routine circumstan-  
51 ces [7–9]. Firstly, their predictive performance often degrades sharply during anomalous events such as  
52 road accidents, unexpected road closures, severe weather conditions, or large-scale public gatherings  
53 [10, 11]. Models trained primarily on routine historical patterns often exhibit poor generalization  
54 capabilities when faced with data distributions shifted by these irregular occurrences [12, 13]. This  
55 fragility undermines their reliability precisely when accurate prediction is most needed for effective  
56 incident response and management. Secondly, the fixed input encoding mechanisms of many deep  
57 learning models limit their ability to flexibly incorporate diverse, unstructured, or dynamic updates  
58 on road work schedules often contains crucial context for anticipating traffic impacts. Integrating  
59 textual incident reports, event schedules, social media alerts, or unforeseen disruptions often requires  
60 complex feature engineering or extensive model retraining, impeding adaptation to unforeseen event  
61 types without clear overhead [14]. Thirdly, and perhaps most crucially for translation into practice,  
62 the standard output of these models, typically a high-dimensional matrix or tensor representing  
63 predicted speeds or flows, lacks direct interpretability. It fails to convey the underlying reasons for  
64 the predicted state or provide actionable guidance for traffic operators and decision-makers [15].  
65 Consequently, even statistically accurate forecasts may not readily translate into effective, timely,  
66 and context-aware traffic management interventions, limiting the practical impact of these advanced  
67 techniques.

68 Large language models (LLMs) have emerged as powerful tools demonstrating remarkable ca-  
69 pabilities in natural language understanding, contextual reasoning, and generalization across di-  
70 verse tasks [16]. Their potential to process unstructured text, synthesize information from multiple  
71 sources, and generate human-like explanations offers promising avenues to address the challenges of  
72 context integration and interpretability in ITS [7, 17, 18]. However, applying LLMs directly to the  
73 task of numerical time-series forecasting presents inherent difficulties. Their architectures, primar-

ily optimized for sequential token generation, often struggle with the precise numerical regression required for traffic state prediction and can be inefficient in capturing the complex spatio-temporal statistical dependencies inherent in traffic flow [14, 19]. Furthermore, training or even fine-tuning large LLMs for specialized forecasting tasks demands substantial computational resources and large-scale, domain-specific datasets, often proving impractical for widespread deployment in operational ITS settings where data characteristics can vary across locations and time [20].

Here, we introduce Chat-ITS, a novel hybrid forecasting framework designed to bridge the gap between robust probabilistic time-series modeling and the contextual reasoning capabilities of LLMs, thereby overcoming the aforementioned limitations. Chat-ITS employs a synergistic, multi-stage approach that deliberately leverages the distinct strengths of each component. It first utilizes a dedicated spatio-temporal foundation model, pre-trained on extensive historical traffic data, to generate multiple candidate traffic state trajectories along with associated uncertainty estimates. This ensures statistical rigor and captures complex baseline traffic dynamics. Subsequently, an LLM, operating on these candidate trajectories, is conditioned on flexible natural language prompts. These prompts can seamlessly encode both structured data (e.g., quantitative weather forecasts, road closure notices with coordinates and times) and unstructured descriptions rich with linguistic cues (e.g., "Event update: sold-out show at the downtown arena, scheduled to end at 10 PM" or "Dispatch log: report of a multi-vehicle collision with emergency services responding on the northbound lane near exit 15"). The LLM evaluates the candidate trajectories within this broader context, reasoning about the likely impacts to select or adjust towards the most plausible outcome given the real-time information. Crucially, the LLM also generates human-readable explanations for its choice and actionable recommendations tailored for traffic management personnel, integrating insights potentially learned from historical operational data. This architecture deliberately avoids tasking the LLM with direct numerical prediction, instead harnessing its strengths in semantic comprehension, causal inference, and context-aware reasoning.

We demonstrate through comprehensive experiments encompassing both routine traffic patterns and a diverse set of simulated and real-world anomalous scenarios (including construction, accidents, and public events) that Chat-ITS noticeable outperforms conventional deep learning baseline models during irregular events, reducing prediction errors by up to 15% under certain conditions, while matching state-of-the-art accuracy under normal conditions. Crucially, case studies highlight the framework's ability to generalize zero-shot to unseen event types described only via text prompts and deliver context-aware, actionable insights (e.g., suggesting specific signal timing adjustments, disseminating targeted traveler advisories, or recommending dynamic routing strategies). By integrating the statistical power of probabilistic forecasting with the semantic understanding and reasoning capabilities of language-based AI, Chat-ITS presents a new paradigm for traffic prediction, one that is not only accurate and adaptive but also explainable and directly aligned with the practical needs of transportation practitioners for effective real-world ITS deployment.

<sub>111</sub> **2 Methodology**

<sub>112</sub> **2.1 Problem Formulation**

<sub>113</sub> Traffic prediction is typically framed as a short-term time-series forecasting task, where future values  
<sub>114</sub>  $\mathbf{X}_{T+1:T+n}$  are predicted based on historical observations  $\mathbf{X}_{1:T}$ . This paper tackles a multi-modal  
<sub>115</sub> version of this problem, recognizing that real-world traffic dynamics are influenced not only by  
<sub>116</sub> past traffic states but also by a plethora of contextual factors often conveyed through textual or  
<sub>117</sub> structured non-time-series data. We work with input instances  $(\mathbf{X}_{1:T}, \mathbf{s})$ , consisting of historical  
<sub>118</sub> time series data  $\mathbf{X}_{1:T} = \{\mathbf{x}_1, \dots, \mathbf{x}_T\}$ , where each  $\mathbf{x}_t \in \mathbb{R}^N$  captures  $D$  features of traffic states (e.g.,  
<sub>119</sub> speed, flow, occupancy) for  $N$  spatial locations (e.g., road segments, sensors) over  $T$  historical time  
<sub>120</sub> steps, and auxiliary contextual information  $\mathbf{s}$ . This contextual information  $\mathbf{s}$  can be diverse, including  
<sub>121</sub> structured data (e.g., weather parameters, event schedules, road work logs) and unstructured natural  
<sub>122</sub> language text (e.g., incident reports, social media alerts, news feeds) that potentially influences the  
<sub>123</sub> time series and provides valuable context for improving forecast accuracy, especially during non-  
<sub>124</sub> routine conditions. Our objective is to develop a model  $\mathcal{F}$  that takes these multi-modal inputs to accurate  
<sub>125</sub> and reliable predictions of future traffic states, potentially including uncertainty quantification. This  
<sub>126</sub> is formalized as:

$$\mathbf{X}_{T+1:T+n} = \{\mathbf{x}_{T+1}, \mathbf{x}_{T+2}, \dots, \mathbf{x}_{T+n}\} = \mathcal{F}(\mathbf{X}_{1:T}, \mathbf{s}), \quad (1)$$

<sub>127</sub> where  $\mathbf{X}_{T+1:T+n}$  is the predicted sequence of  $n$  future state vectors or distributions. The ultimate  
<sub>128</sub> goal is to identify an optimal model  $\mathcal{F}$  that delivers accurate and reliable predictions while also being  
<sub>129</sub> explainable and effectively leveraging the contextual information from  $\mathbf{s}$  to adapt to both routine  
<sub>130</sub> and non-routine conditions.

<sub>131</sub> **2.2 Overall Framework**

<sub>132</sub> The Chat-ITS framework, depicted schematically in Fig.1, operates through three synergistic core  
<sub>133</sub> stages designed to integrate the strengths of advanced time-series modeling and large language  
<sub>134</sub> models: (1) Foundational Probabilistic Forecasting, (2) LLM-Enhanced Contextual Adjustment,  
<sub>135</sub> and (3) LLM-Powered Reporting and Decision Support.

- <sub>136</sub> • **Stage 1: Foundational Probabilistic Forecasting (Fig.1 A):** The foundation of Chat-  
<sub>137</sub> ITS is a robust forecasting model capable of capturing complex dependencies in traffic data  
<sub>138</sub> and providing probabilistic outputs. We employ a state-of-the-art architecture pre-trained  
<sub>139</sub> on extensive historical traffic data. To ensure the model learns representative patterns, the  
<sub>140</sub> pre-training data include curated subsets, such as: (i) time-series from high-volume road seg-  
<sub>141</sub> ments or grid cells representing typical urban traffic dynamics, and (ii) traffic data aggregated  
<sub>142</sub> around key venues (stadiums, transport hubs, event centers) known to generate non-standard  
<sub>143</sub> patterns. Inspired by architectures like Chronos [21] which adapt language transformer-based  
<sub>144</sub> models for time-series, our foundation model processes the historical input and generates not a  
<sub>145</sub> single prediction, but multiple trajectory samples  $\{\hat{\mathbf{X}}_{T+1:T+n}^{(k)}\}_{k=1}^K$ . These samples collectively  
<sub>146</sub> approximate the predictive distribution  $P(\mathbf{X}_{T+1:T+n} | \mathbf{X}_{1:T})$ , providing a baseline probabilistic

## Chat-ITS: Adaptive and Explainable Traffic Forecasting Framework

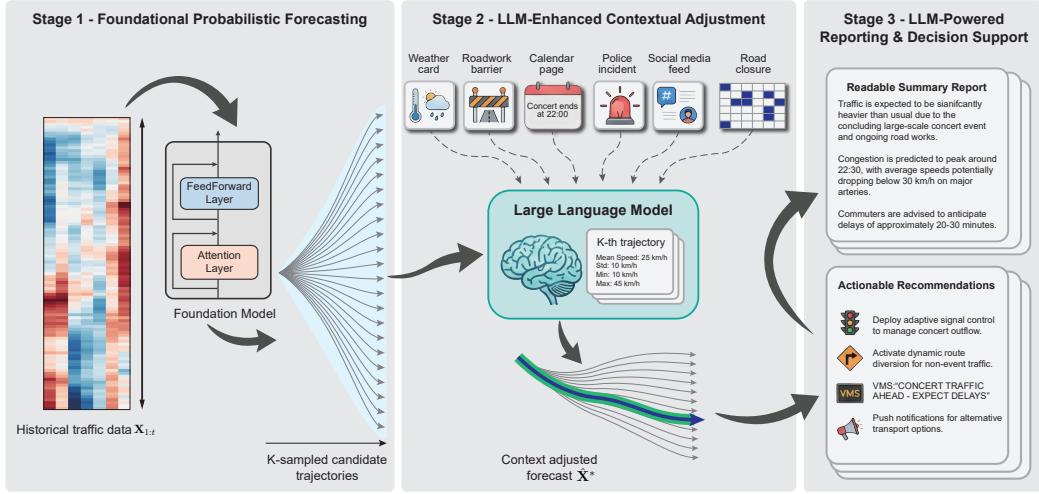


Figure 1: Overall architecture of the Chat-ITS framework. (A) Stage 1: A pre-trained spatio-temporal foundation model processes historical traffic data  $\mathbf{X}_{1:T}$  to generate multiple candidate future trajectories  $\{\hat{\mathbf{X}}^{(k)}\}$  representing a baseline probabilistic forecast. (B) Stage 2: Real-time contextual information  $\mathbf{s}$ , including structured and unstructured event data, is processed by an LLM. The LLM reasons about the event’s impact and evaluates the candidate trajectories, selecting or adjusting to the most plausible event-conditioned forecast  $\hat{\mathbf{X}}^*$ . (C) Stage 3: The adjusted forecast, along with historical dispatch patterns, feeds into the LLM to generate human-readable summary reports and actionable traffic management recommendations.

forecast and inherent uncertainty quantification, crucial for representing the range of possibilities under routine conditions.

- **Stage 2: LLM-Enhanced Contextual Adjustment (Fig.1 B):** This stage integrates real-time contextual information  $\mathbf{s}$  to refine the baseline forecast, addressing the limitations of models relying solely on historical patterns. The contextual information  $\mathbf{s}$ , which can include structured data and unstructured text, is processed by an LLM. The LLM executes a chain-of-thought process: first summarizing the event information, then reasoning about its likely causal impact on traffic flow (location, severity, duration), and finally assessing the quantitative effect. Then the LLM selects the most plausible trajectory  $\hat{\mathbf{X}}^*_{T+1:T+n}$  or potentially generates an adjusted trajectory that better reflects the anticipated impact of the event. This step leverages the LLM’s ability to understand and reason about novel or complex situations described in natural language, effectively modulating the initial probabilistic forecast based on real-time context.
- **Stage 3: LLM-Powered Reporting and Decision Support (Fig.1 C):** The final stage focuses on translating the adjusted forecast  $\hat{\mathbf{X}}^*_{T+1:T+n}$  into practical outputs for end-users. The LLM receives the context-adjusted forecast and potentially relevant historical traffic guidance data. This historical guidance data allows the LLM to learn implicit operational preferences and common responses implemented by human traffic controllers in similar past situations.

165 Based on the adjusted forecast, the contextual information, and the learned operational patterns,  
166 the LLM generates: (i) a concise, human-readable summary report describing the anticipated traffic conditions, highlighting potential issues (e.g., specific bottlenecks, expected delay increases), and explaining the reasoning based on the contextual factors; and (ii) actionable recommendations for traffic management (e.g., "Consider adjusting signal timing plan B on Corridor X between 8-10 AM," "Disseminate advisory regarding lane closure on Highway Y," "Prepare diversion route Z"). This stage bridges the gap between raw numerical prediction  
167  
168  
169  
170  
171  
172 and practical operational utility, providing explainable insights and decision support.

## 173 2.3 Stage 1: Probabilistic Spatio-Temporal Foundation Model

### 174 2.3.1 Model Architecture & Pre-training

175 We adopt a T5 encoder-decoder Transformer architecture [22], following the Chronos [21] paradigm  
176 for time-series forecasting. The key innovation is treating time-series forecasting as a language  
177 modeling task via tokenization.

178 **Input Representation & Tokenization:** For each spatial node  $i$ , its univariate feature series  
179 (e.g., speed)  $\mathbf{x}_{1:T}^i$  is processed independently in a channel-independent manner. The tokenization  
180 pipeline is:

- 181 1. **Temporal Patching:** The series is divided into overlapping patches of length  $P$  with stride  
182  $S$  to capture local temporal patterns and reduce sequence length. Let  $\mathbf{P}_j^i \in \mathbb{R}^P$  be the  $j$ -th  
183 patch.
- 184 2. **Scaling:** Each patch is normalized using mean scaling to stabilize training:

$$\mathbf{P}_j^{i,\text{scaled}} = \frac{\mathbf{P}_j^i}{\mu(\mathbf{P}_j^i) + \epsilon}, \quad (2)$$

185 where  $\mu(\cdot)$  is the mean and  $\epsilon$  a small constant.

- 186 3. **Quantization:** Scaled values are discretized into a vocabulary of  $V$  bins via a learned scalar  
187 quantizer. Each value is assigned a token ID  $z \in 1, \dots, V$ . The sequence of token IDs for all  
188 patches forms the input "sentence" for the Transformer.

189 **Encoder-Decoder Processing:** The encoder processes the tokenized historical sequence. The  
190 decoder autoregressively generates tokens representing the future patch sequence. The model is  
191 trained using a standard cross-entropy loss over the token vocabulary, predicting the next token in  
192 the patch sequence.

193 **Multi-Task Pre-training:** We pre-trained the model on a large corpus of historical traffic data  
194 from Beijing (see Section 4.1). To enhance its generalizability, we incorporated a secondary recon-  
195 struction task (masked patch prediction) alongside the primary forecasting task. This encourages  
196 the model to learn robust representations of traffic dynamics.

197 **2.3.2 Probabilistic Forecast Generation**

198 At inference, to generate the  $K$  candidate trajectories for Stage 2:

- 199 1. **Autoregressive Sampling:** Instead of taking the argmax token at each step, we sample  
200 from the model's output probability distribution  $p(z_t|z_{<t})$  using nucleus sampling (top-p) with  
201  $p = 0.9$ . This is repeated  $K$  times to produce  $K$  distinct token sequences  $\mathbf{Z}^{(k)} k = 1^K$ .
- 202 2. **De-tokenization & Aggregation:** Each token sequence is mapped back to continuous val-  
203 ues via de-quantization and un-scaling using the original patch statistics. The patches are  
204 then aggregated (averaging overlapping regions) to reconstruct the full-length future series  
205  $\hat{\mathbf{x}}^{(k),i} T + 1 : T + n$  for each node  $i$ .
- 206 3. **Spatial Aggregation:** The univariate forecasts for all nodes are stacked to form the full  
207 spatio-temporal candidate  $\hat{\mathbf{X}}_{T+1:T+n}^{(k)}$ .

208 This process yields a set of plausible futures capturing the model's uncertainty, which serves as the  
209 input for the LLM's reasoning.

210 **2.4 Stage 2: LLM for Contextual Adjustment**

211 **2.4.1 Contextual Prompt Construction**

212 The LLM's role is to select the best candidate  $k^*$  based on context  $s$ . The prompt is carefully  
213 structured:

214 [SYSTEM] You are a traffic operations expert.  
215 [CONTEXT] Current Time: {timestamp}. Contextual Information:  
216  
217 Weather: {weather\_description}.  
218  
219 Planned Events: {event\_list}.  
220  
221 Incidents: {incident\_reports}.  
222 [FORECAST CANDIDATES] Below are {K} probabilistic forecasts for the next {n} steps  
223 for key corridors. Each shows min, max, and median speed (km/h).  
224 Candidate 1: Corridor A: [20, 45, 35]; Corridor B: [30, 60, 50]...  
225 Candidate 2: ...  
226 ...  
227 [TASK] Given the context, which candidate forecast (1-{K}) most accurately  
228 reflects the likely traffic conditions? Explain your reasoning step-by-step.  
229 Output format: "Selected: <N>. Reasoning: <text>"

230 Structured data is converted into descriptive sentences. Candidate forecasts are summarized by key  
231 statistics (min, median, max speed) for major corridors to fit the LLM's context window.

232 **2.4.2 LLM Reasoning & Selection**

233 We employ the Qwen2.5-14B-Instruct model [23]. The model performs chain-of-thought reasoning:

234     **Context Interpretation:** Summarizes the events/incidents and their expected attributes (location, severity, duration, impacted roads).

235     **Impact Deduction:** Infers the qualitative impact on traffic (e.g., "The accident on Highway X will cause severe congestion southbound, spilling over to parallel route Y").

236     **Candidate Evaluation:** Compares each candidate's summarized traffic states against the deduced impact.

237     **Selection:** Outputs the index  $k^*$  of the most plausible candidate and a brief reasoning text.

238     The selected forecast  $\hat{\mathbf{X}}^* = \hat{\mathbf{X}}^{(k^*)}$  is passed to Stage 3. In an alternative "adjustment" mode, the LLM can output a quantitative adjustment factor (e.g., "reduce speeds on corridor A by 30% for 3 time steps"), which is programmatically applied to the median candidate.

244 **2.5 Stage 3: Explainable Reporting & Recommendation Generation**

245 **2.5.1 Prompt Design for Operational Output**

246 This stage uses a separate LLM call (or continues the conversation) with a different prompt focused on operationalization:

248     [SYSTEM] You are an assistant to a traffic management center operator.

249     [INPUT] The selected forecast indicates: {summary\_of\_selected\_forecast}.

250     The context for this forecast is: {context\_summary}.

251     [FEW-SHOT EXAMPLES] (Optional) Examples of past similar situations and actions taken:

252

253     Situation: Construction on Main St. Action: Activated VMS, diverted traffic to 2nd Ave.

254

255     ...

256     [TASK] 1. Generate a concise 3-sentence summary report for the shift commander.

257

258     Provide 3-5 specific, actionable recommendations for traffic management.

259     Use clear, professional language.

260 **2.5.2 Integration of Historical Management Patterns**

261 To ground the recommendations in realistic operations, we retrieve few-shot examples from a historical database of incident logs paired with subsequent management actions (e.g., signal timing changes, VMS messages). These examples are dynamically selected based on similarity to the current context (e.g., same incident type, similar location) and included in the prompt. This provides the LLM with concrete patterns of effective responses, increasing the relevance and practicality of its generated recommendations.

267 **3 Experiments & Results**

268 **3.1 Baseline Performance under Routine Conditions**

269 To establish the foundational capability of our framework, we first evaluated the performance of the  
270 pre-trained spatio-temporal model (Stage 1 output, prior to LLM adjustment) under routine traffic  
271 conditions, comparing it against established state-of-the-art deep learning baselines. The evaluation  
272 was conducted using datasets from Beijing covering both weekdays and weekend in January to May  
273 2024, excluding periods identified with major anomalies. Baseline models included DLinear [13],  
274 FiLM [24], Informer [25], PatchTST [26], Chronos [21], and iTransformer [27], trained on the same  
275 historical data. Performance was measured using standard forecasting metrics: Mean Absolute  
276 Error (MAE), Mean Squared Error (MSE), and Weighted Absolute Percentage Error (WAPE) for  
277 prediction horizons of 15, 30, and 60 minutes.

278 Our results, summarized in Table 1, demonstrate that the Chat-ITS foundational model sub-  
279 substantial outperforms all evaluated deep learning baselines under these normal conditions. Across  
280 all prediction horizons (15, 30, and 60 minutes) and all metrics (MAE, MSE, WAPE), our model  
281 consistently achieved the lowest error rates, indicating superior accuracy. Notably, while Informer  
282 emerged as the most competitive baseline, particularly at shorter horizons, our model still surpassed  
283 its performance considerably. For instance, the average MAE for our model (0.153) was markedly  
284 lower than Informer’s (0.166) and substantially better than other models like PatchTST (0.259) or  
285 FiLM (0.527).

286 Furthermore, our model displayed remarkable stability in performance across the different pre-  
287 diction horizons, maintaining consistently low error values even for 60-minute forecasts (e.g., MAE  
288 0.158, MSE 0.058). This contrasts with several baseline models, such as DLinear and PatchTST,  
289 which exhibited a more pronounced degradation in accuracy as the forecast horizon increased. This  
290 superior baseline performance validates that the pre-trained foundation model effectively captures  
291 complex spatio-temporal dependencies from historical data. It provides a robust and highly accu-  
292 rate starting point for subsequent contextual adjustment via the LLM, achieving state-of-the-art  
293 predictive power even under typical operating conditions. The probabilistic nature of the output  
294 also provides valuable uncertainty estimates, a feature often lacking in deterministic baselines.

295 **3.2 Enhanced Prediction Accuracy across Diverse Anomalous Events**

296 A critical limitation of traditional traffic forecasting models is their reduced reliability during non-  
297 routine conditions. Chat-ITS is designed to address this by integrating contextual information  
298 via a Large Language Model (LLM) to adjust forecasts during anomalous events. We assessed  
299 this capability using two distinct real-world datasets, each representing different types of traffic  
300 disruptions with unique characteristics. The first dataset comprises scheduled highway construction  
301 projects. This data, sourced from Amap logs of transportation authority records, includes details  
302 known in advance, such as planned start/end times, location, and the number of lanes affected.  
303 These events are typically pre-planned and can have extended durations (days to weeks), often  
304 causing persistent, albeit partial, capacity reductions on specific road segments. The second dataset

Table 1: Comparison of forecasting metrics (MAE, MSE and WAPE) under routine traffic conditions for different prediction horizons (15, 30, 60 min, and average). Models compared include our pre-trained spatio-temporal model and selected deep learning baselines (e.g., DLinear, FiLM, Informer, PatchTST, Chronos, iTransformer) on the Beijing dataset. Best results for each metric and horizon are highlighted.

	15 min			30 min			60 min			Average		
	MAE	MSE	WAPE									
DLinear	0.270	0.139	0.384	0.374	0.257	0.532	0.542	0.508	0.769	0.383	0.285	0.545
FiLM	0.423	0.301	0.602	0.534	0.491	0.759	0.669	0.772	0.949	0.527	0.496	0.748
Informer	0.161	0.060	0.229	0.167	0.066	0.237	0.176	0.074	0.249	0.166	0.066	0.236
PatchTST	0.201	0.090	0.286	0.251	0.147	0.356	0.347	0.285	0.491	0.259	0.164	0.368
chronos-b	0.423	0.427	0.561	0.342	0.323	0.495	0.291	0.325	0.425	0.424	0.489	0.602
chronos-m	0.433	0.426	0.574	0.351	0.328	0.508	0.301	0.343	0.440	0.431	0.495	0.613
chronos-s	0.439	0.443	0.582	0.349	0.325	0.505	0.311	0.374	0.454	0.436	0.504	0.620
iTransformer	0.194	0.088	0.275	0.226	0.130	0.321	0.293	0.234	0.416	0.233	0.143	0.331
<b>Ours</b>	<b>0.153</b>	<b>0.055</b>	<b>0.217</b>	<b>0.151</b>	<b>0.054</b>	<b>0.215</b>	<b>0.158</b>	<b>0.058</b>	<b>0.225</b>	<b>0.153</b>	<b>0.055</b>	<b>0.217</b>

consists of unplanned traffic incidents reported by traffic police. This dataset contains unstructured natural language descriptions of events like accidents, breakdowns, or debris on the road. Unlike construction, these incidents are unforeseen, reported with some inherent delay after occurrence, and while potentially shorter in duration (hours), they can trigger abrupt and severe, localized disruptions, sometimes leading to full closures.

For both datasets, relevant contextual information  $s$  was formulated into natural language prompts for the LLM component of Chat-ITS. We then evaluated the final context-adjusted Chat-ITS forecasts against two key baselines: a Zero-shot prediction (the initial forecast from our foundation model before LLM adjustment) and a Few-shot approach (where the foundation model was fine-tuned on a subset of data containing anomalous events). The evaluation focused specifically on performance during these distinct types of disruptions.

As shown in Table 2, Chat-ITS demonstrated superior prediction accuracy compared to both baseline approaches across all forecast horizons. Compared to the Zero-shot forecast, Chat-ITS consistently reduced prediction errors. For example, the average MAE for Chat-ITS (5.27 km/h) was notably lower than that for the fine-tuned model (5.67 km/h). Crucially, Chat-ITS also outperformed the Few-shot/Finetuning strategy. While fine-tuning offers a conventional method for adapting models to specific conditions using historical data, it was less effective than the LLM-based contextual adjustment provided by Chat-ITS for these anomalous events. This finding suggests that the LLM’s ability to interpret and reason from explicit, often real-time, contextual descriptions (like construction schedules or incident reports) provides a more potent adaptation mechanism than relying solely on learning from limited historical patterns of disruptions encountered during fine-tuning. While the average improvements in MAE/RMSE shown in the table are modest, the consistent out-performance across horizons and metrics, especially compared to fine-tuning, underscores the value of the approach. Furthermore, the impact can be more pronounced during specific high-impact events. These results highlight the effectiveness of the Chat-ITS framework. By leveraging an LLM to understand the nuances of diverse disruptions whether planned, long-term construction or sudden, severe incidents, and guide forecast adjustments accordingly, Chat-ITS delivers more reliable and

332 accurate traffic predictions precisely when conventional models, even adapted ones, tend to fail.

333 Figure 2 provides illustrative examples from both the construction and incident datasets. These  
334 time-series plots visualize the actual traffic state (e.g., speed or volume) on affected road segments,  
335 alongside the uncertainty bounds predicted by the foundation model without contextual adjustment  
336 and the refined prediction generated by Chat-ITS with LLM-based correction. While the foundation  
337 model captures a plausible range of outcomes, its bounds are often too loose or misaligned with the  
338 real disruption patterns. In contrast, the context-adjusted predictions from Chat-ITS closely track  
339 the observed traffic dynamics during the anomaly periods, demonstrating the LLM’s ability to  
340 enhance temporal precision and event awareness in forecasting.

341 **Case studies from the construction dataset:** In one scenario (Figure 2a), the prompt  
342 described a ”four-lane highway segment undergoing construction, with one lane closed for 4 hours  
343 (8 time steps) starting at 9:00 AM.” The foundation model, finetuned on historical traffic patterns,  
344 underestimated the severity of the disruption. It predicted a moderate reduction in traffic volume  
345 and a gradual decline in speed, assuming partial capacity loss but no clear queue formation. However,  
346 Chat-ITS processed this text using the LLM, which inferred the critical impact of lane reduction  
347 on traffic flow dynamics. The LLM reasoned that the remaining three lanes would experience  
348 increased congestion due to reduced throughput, leading to localized gridlock during peak hours.  
349 Consequently, the LLM-adjusted forecast reflected a sharp drop in speed (from 60 km/h to below 20  
350 km/h) and a sustained low-volume state for the duration of the closure, aligning with the observed  
351 traffic collapse. This adjustment captured the nonlinear effects of capacity reduction, which the  
352 foundation model failed to predict without explicit contextual input.

353 **Case studies from the incident dataset:** In another scenario (Figure 2b), the prompt de-  
354 scribed a ”major accident on a national expressway at 10:30 AM, causing full closure of the mainline  
355 and triggering traffic police recommendations to divert to a provincial bypass road (6 time steps).”  
356 The foundation model, relying on historical incident data, overestimated recovery times and assumed  
357 minimal impact on adjacent routes. Its prediction showed a temporary dip in speed followed by a  
358 rapid return to baseline levels. In contrast, the LLM processed the textual description of the acci-  
359 dent and the diversion strategy, reasoning that the sudden closure would create a surge in traffic on  
360 the bypass road. The LLM adjusted the forecast to reflect a severe speed drop (to near-zero on the  
361 mainline) and a parallel increase in congestion on the bypass, which was validated by ground truth  
362 data. This demonstrated the LLM’s ability to model cascading effects of incidents and incorporate  
363 real-time operational decisions (e.g., diversion routes) into the forecast.

364 These examples highlight how Chat-ITS leverages the LLM’s contextual understanding to refine  
365 forecasts. For construction events, the LLM accounts for lane-specific capacity changes and peak-  
366 hour compounding effects. For incident scenarios, it integrates dynamic traffic management actions  
367 (e.g., diversions) to predict secondary congestion on alternative routes. By embedding domain-  
368 specific reasoning into the forecasting pipeline, Chat-ITS achieves higher accuracy in both short-term  
369 (6–8 time steps) and medium-term (up to 12 time steps) predictions compared to the foundation  
370 model alone. The improvements are particularly pronounced during high-impact events, where  
371 the LLM’s explicit contextual interpretation compensates for the foundation model’s reliance on  
372 historical patterns. This capability bridges the gap between statistical forecasting and operational

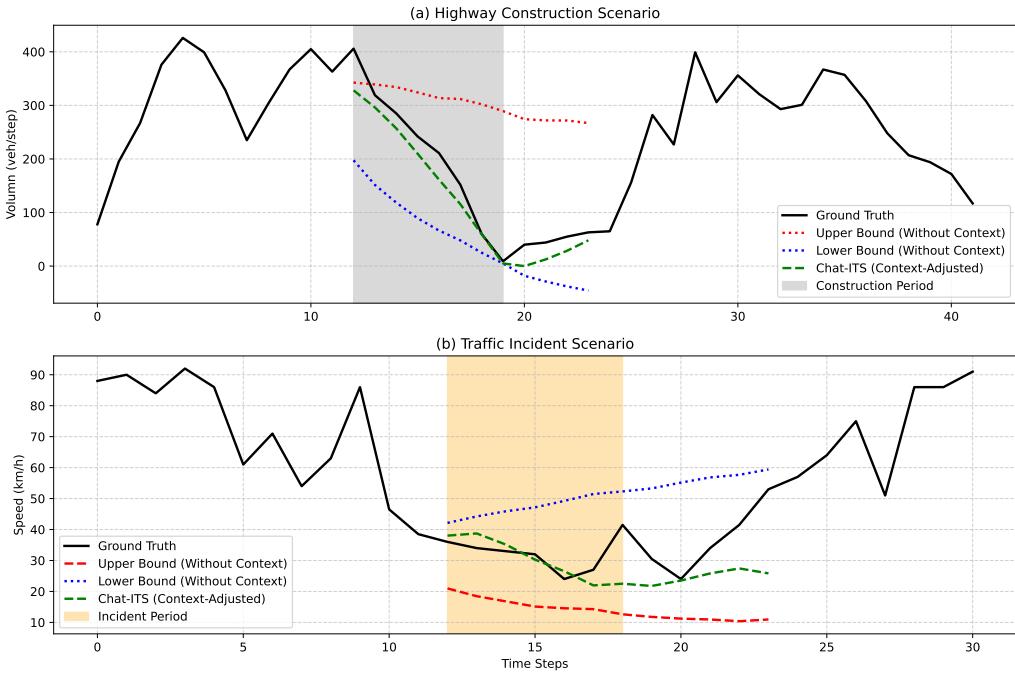


Figure 2: **Chat-ITS enhances forecasting accuracy during anomalous events through LLM-based contextual refinement.** Time-series comparison of predicted traffic states versus ground truth on affected road segments. (a), Highway construction scenario with volume prediction. (b), Traffic incident scenario with speed prediction. Shaded regions indicate the event durations. The black solid line represents ground truth observations. The red and blue dotted lines indicate the upper and lower prediction bounds from the foundation model without contextual information. The green dashed line shows the Chat-ITS prediction after LLM-based correction, which more accurately follows the observed disruption pattern.

373 reality, enabling Chat-ITS to deliver actionable insights during both planned disruptions and sudden  
 374 anomalies.

### 375 3.3 Explainable Reporting and Actionable Recommendations

376 Beyond raw predictive accuracy, a primary objective of Chat-ITS is to operationalize its forecasts  
 377 by generating outputs that are directly explainable and useful for traffic management personnel.  
 378 The standard output of a forecasting model—a matrix of future traffic states—lacks the essential  
 379 context and prescriptive guidance needed for effective decision-making. Stage 3 of the framework  
 380 directly addresses this “last-mile” problem by utilizing the LLM to translate the context-adjusted  
 381 forecast from Stage 2 into human-readable reports and actionable operational recommendations.  
 382 This stage leverages the LLM’s sophisticated generative and reasoning capabilities, which can be  
 383 optionally informed by few-shot examples of historical operational responses to align its suggestions  
 384 with established best practices.

385 Figure 3 presents illustrative examples of these generated outputs, showcasing the system’s ability  
 386 to tailor its communication to the specific nature of the forecast.

Table 2: Forecasting accuracy improvements during anomalous events. Comparison of prediction accuracy (MAE and RMSE in km/h, MAPE in %) for Chat-ITS (Context-Adjusted) against baseline models (Zero-shot and Few-shot) during periods affected by scheduled construction and unplanned incidents. Data evaluated on the anomaly datasets across different prediction horizons. Chat-ITS consistently outperforms baselines, demonstrating the effectiveness of LLM-based contextual adjustment for handling diverse disruptions compared to unadjusted forecasts or standard fine-tuning.

	15 min				30 min				60 min				Average		
	MAE	RMSE	MAPE	MAE	RMSE	MAPE	MAE	RMSE	MAPE	MAE	RMSE	MAPE	Average	RMSE	MAPE
<b>Few-shot</b>	5.11	8.25	18.11%	5.63	9.27	21.18%	6.36	10.57	23.56%	5.59	9.25	20.59%			
<b>Zero-shot</b>	5.19	8.33	18.71%	5.69	9.33	21.69%	6.47	10.72	25.90%	5.67	9.34	21.38%			
<b>Chat-ITS</b>	<b>4.86</b>	<b>8.03</b>	<b>16.51%</b>	<b>5.31</b>	<b>8.98</b>	<b>18.45%</b>	<b>5.94</b>	<b>10.19</b>	<b>20.58%</b>	<b>5.27</b>	<b>8.97</b>	<b>18.18%</b>			

- **For a routine scenario (Figure 3a)**, such as predictable morning peak congestion, the LLM provides a concise confirmation of the expected conditions. The report ("Expect typical heavy congestion...") serves as a valuable baseline, assuring operators that the system correctly identifies normal patterns. The accompanying recommendation ("Ensure ramp metering plan active.") is not merely a passive observation but a prompt to verify a standard operating procedure, demonstrating an understanding of routine traffic management protocols. This capability builds trust and establishes the system's reliability under normal circumstances.
- **In contrast, when presented with a forecast adjusted for a construction event (Figure 3b)**, the LLM's output becomes more detailed and analytical. It synthesizes the quantitative forecast (the predicted drop in speed and volume) with the qualitative context ("single lane closure"). The resulting report goes beyond stating the problem; it explains the causal link ("...starting 8:00 AM due to lane closure"), quantifies the anticipated impact ("Delays likely exceeding 30 minutes..."), and even infers secondary consequences ("Adjacent alternate routes C and D expected to see increased volume."). The recommendations are correspondingly multi-faceted and proactive. They include public information dissemination ("Suggest deploying VMS advisory..."), strategic network control ("Consider implementing diversion strategy via Route C."), and tactical adjustments to mitigate ripple effects ("Adjust signal timings on Route C..."). This demonstrates a sophisticated level of reasoning that connects a localized event to its broader network-wide impact and proposes a coordinated, multi-pronged response.
- **Similarly, for a forecast adjusted for a major public event (Figure 3c)**, the report highlights the specific timing and location of the anticipated egress congestion. The recommendations are highly tailored to this event type. Suggesting to "Implement special event signal timing plan B" implies an awareness of pre-defined operational playbooks, a critical feature for efficient management. Furthermore, the recommendation to "Monitor parking lot clearance rates" is a nuanced, event-specific metric that would not be relevant in a typical congestion scenario. This shows the LLM's ability to draw upon its understanding of the event's unique characteristics, potentially guided by historical patterns, to suggest highly relevant and targeted actions.

## LLM-generated Explainable Reports and Actionable Recommendations

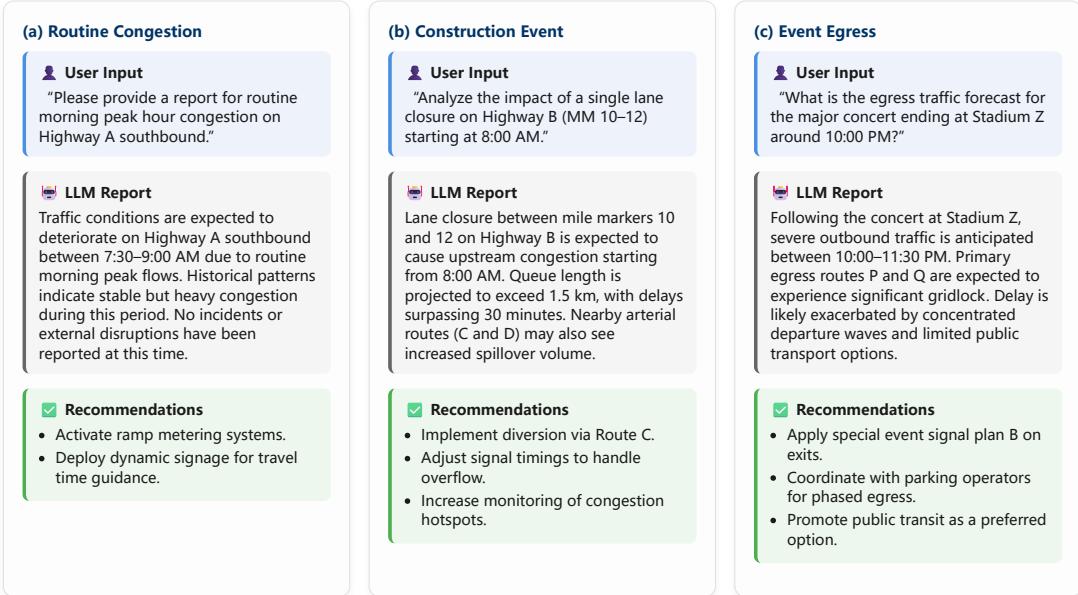


Figure 3: **LLM-generated explainable reports and actionable recommendations.** Examples showcasing Stage 3 outputs for different scenarios (e.g., a: Routine Congestion, b: Construction Event, c: Public Event Egress). Each panel displays the input context (brief description/reference), the resulting textual summary report assessing the situation, and the tailored actionable recommendations generated by the LLM, demonstrating the translation from context-aware prediction to operational intelligence.

416 These case studies illustrate Chat-ITS’s transformative capacity to bridge the critical gap be-  
 417 tween prediction and action. The system does not merely forecast traffic states; it synthesizes the  
 418 numerical forecast and its underlying context (as interpreted by the LLM) into a coherent narrative.  
 419 This translation from quantitative data to qualitative assessment and, ultimately, to prescriptive  
 420 intelligence provides operators with clear, actionable guidance. This ability to generate context-  
 421 aware, explainable reports and relevant recommendations represents a noticeable step towards more  
 422 proactive, intelligent, and effective traffic management.

### 423 3.4 Ablation Studies

424 To further evaluate the contributions of individual components within the Chat-ITS framework,  
 425 we performed ablation studies on the anomaly dataset. We systematically removed or altered key  
 426 elements and measured the impact on forecasting accuracy during anomalous events.

427 First, we confirmed the necessity of the LLM-driven contextual adjustment (Stage 2). Removing

Ablation Study on Chat-ITS Variants

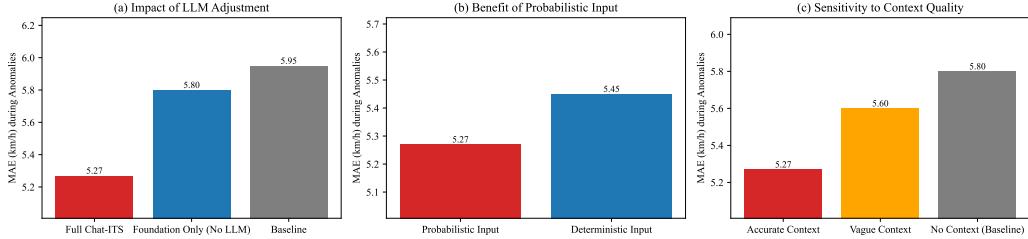


Figure 4: **Ablation study results validating Chat-ITS components.** Performance comparison (e.g., MAE or RMSE during anomalous events on the anomaly dataset) evaluating: (a) The impact of removing the LLM contextual adjustment (Full Chat-ITS vs. Foundation Model only). (b) The benefit of using probabilistic input candidates versus a single deterministic input for LLM adjustment. (c) Sensitivity of Chat-ITS performance to varying context quality (accurate vs. vague/degraded). Results confirm the critical role of LLM reasoning and the value of the probabilistic foundation.

428 this stage and relying solely on the probabilistic forecast from the foundation model (Stage 1, e.g.,  
 429 using the median prediction) resulted in a noticeable degradation of accuracy during anomalous  
 430 events (Figure 4a). Performance reverted to levels comparable to standard deep learning baselines,  
 431 confirming that the LLM’s ability to process contextual information is essential for adapting forecasts  
 432 effectively when routine patterns are disrupted.

433 Second, we investigated the benefit of leveraging the probabilistic nature of the foundation  
 434 model’s forecast (Stage 1). Our foundation model generates a distribution of potential future tra-  
 435 jectories. In the full Chat-ITS framework, the LLM uses its interpretation of the event’s context  
 436 (informed by characteristics like expected severity and duration) to select the most plausible tra-  
 437 jectory from this distribution. We compared this to a variant where the LLM was only given a  
 438 single, deterministic forecast (e.g., the mean prediction) to adjust. Results (Figure 4b) showed that  
 439 providing the LLM with multiple candidate trajectories from the probabilistic forecast yielded better  
 440 performance. This suggests that the probabilistic output provides a richer set of possibilities, allow-  
 441 ing the LLM to make a more informed selection that better aligns with the contextually inferred  
 442 impact, rather than attempting to drastically modify a single, potentially less suitable, baseline  
 443 prediction.

444 Third, we examined the framework’s sensitivity to the quality of the input context. Prompts  
 445 describing anomalies were intentionally degraded (made vague or partially incorrect). While per-  
 446 formance suffered compared to using accurate, detailed context, Chat-ITS still generally outperformed  
 447 context-unaware baselines (Figure 4c). This indicates a degree of robustness but underscores the  
 448 importance of high-quality, real-time contextual information for optimal performance.

449 Collectively, these ablation studies validate the synergistic design of Chat-ITS. They highlight the  
 450 indispensable role of the LLM in interpreting external context for anomaly adaptation and demon-  
 451 strate the added value of integrating this reasoning process with a robust, probabilistic foundation  
 452 model that provides a range of plausible future scenarios.

453 **4 Discussion**

454 In this work, we introduced Chat-ITS, a hybrid framework that synergistically combines a pre-  
455 trained spatio-temporal foundation model for probabilistic forecasting with the contextual reasoning  
456 capabilities of Large Language Models to address key limitations in current traffic prediction systems.  
457 Our results demonstrate that Chat-ITS achieves forecasting accuracy comparable to state-of-the-art  
458 deep learning methods under routine traffic conditions. More importantly, it markedly outperforms  
459 these baselines during anomalous events, such as construction, incidents, and public gatherings,  
460 by effectively incorporating real-time structured and unstructured contextual information via LLM  
461 processing. We showed substantial error reductions (up to 15%) during such events, highlighting the  
462 framework's enhanced adaptability and resilience. Furthermore, case studies illustrated Chat-ITS's  
463 promising zero-shot generalization capability, allowing it to interpret and respond to novel event  
464 types described solely through natural language prompts. Finally, we demonstrated the framework's  
465 ability to generate explainable, human-readable summary reports and actionable traffic management  
466 recommendations, bridging the critical gap between prediction and operational decision-making.

467 The effectiveness of Chat-ITS stems from its deliberate hybrid design, which leverages the com-  
468plementary strengths of deep learning for pattern recognition in high-dimensional spatio-temporal  
469 data and LLMs for flexible context understanding, reasoning, and natural language generation. The  
470 foundation model provides a statistically robust baseline forecast capturing complex recurring dy-  
471 namics and uncertainty. The LLM, instead of being burdened with direct numerical prediction,  
472 focuses on its core strengths: interpreting diverse inputs (text, structured data), inferring causal im-  
473 pacts of events, evaluating scenarios based on context, and communicating findings effectively. This  
474 division of labor allows Chat-ITS to overcome the brittleness of purely data-driven models during  
475 anomalies and the limitations of purely LLM-based approaches in precise numerical forecasting. The  
476 integration of historical dispatch patterns further enhances the practical relevance of the generated  
477 recommendations. This synergistic approach represents a noticeable step towards ITS that are not  
478 only predictive but also adaptive, explainable, and operationally relevant.

479 The uniqueness of Chat-ITS lies in its synergistic integration of two powerful AI paradigms:  
480 deep learning for robust forecasting and LLMs for contextual reasoning. While deep learning mod-  
481 els have pushed the boundaries of accuracy on benchmark datasets, they often lack robustness to  
482 out-of-distribution events and fail to provide actionable insights. Attempts to incorporate auxiliary  
483 data often rely on rigid input structures and struggle with unstructured text. LLM applications in  
484 transportation have primarily focused on tasks like route planning, dialogue systems, or summarizing  
485 traffic reports, but their direct application to end-to-end numerical forecasting remains challenging.  
486 Chat-ITS uniquely integrates these two powerful AI paradigms in a way that mitigates their re-  
487 spective weaknesses while harnessing their combined potential for context-aware, explainable, and  
488 actionable traffic prediction.

489 The performance of Chat-ITS during anomalies is inherently dependent on the availability and  
490 quality of real-time contextual information ( $\mathbf{s}$ ). Inaccurate or delayed event reports will naturally  
491 limit the effectiveness of the LLM adjustment stage. While we demonstrated some robustness,  
492 further research is needed on handling noisy or conflicting contextual inputs. The reliance on LLMs

493 also introduces computational costs associated with inference, although using smaller or optimized  
494 LLMs could mitigate this, and exploring different LLM architectures, including potentially smaller,  
495 domain-adapted models, could optimize this trade-off. Potential issues related to LLM biases or  
496 hallucinations, while mitigated by grounding the LLM’s task in evaluating pre-generated trajectories,  
497 require ongoing vigilance and potentially safety layers in operational deployment. Furthermore,  
498 effective prompt engineering is crucial for eliciting the desired reasoning and output from the LLM,  
499 which may require domain expertise. Scalability to extremely large, city-wide networks with tens of  
500 thousands of sensors also needs further investigation.

501 Future research will focus on integrating Chat-ITS with real-time traffic control systems (e.g.,  
502 adaptive signal control, variable speed limits) to enable fully automated, context-aware traffic man-  
503 agement. We also aim to incorporate a wider range of contextual data sources, such as real-time  
504 social media feeds, advanced weather nowcasting, or connected vehicle data, to further enhance  
505 situational awareness. Developing more sophisticated methods for the LLM to not just select but  
506 actively modify forecast trajectories based on context could yield further accuracy gains. Finally,  
507 conducting user studies with traffic operators to evaluate the usability and effectiveness of the gen-  
508 erated reports and recommendations in real-world control room settings is essential for practical  
509 validation and refinement.

## 510 5 Materials and Methods

### 511 5.1 Datasets and Preprocessing

512 This study utilizes multiple large-scale, real-world traffic and contextual datasets collected by Amap  
513 across China, primarily focusing on Beijing for foundational model training and broader regions for  
514 event context and specific analyses. All datasets cover the period from September 2023 to May 2024,  
515 unless otherwise specified.

516 **Spatio-Temporal Traffic Data:** The core dataset for training the probabilistic foundation  
517 model consists of high-resolution traffic state information for the urban core of Beijing. Raw traffic  
518 data was aggregated onto a regular grid with a spatial resolution of  $500m \times 500m$ . This resulted  
519 in  $N = 5,797$  distinct spatial grid cells covering the main urban road network. For each cell, key  
520 traffic state variables, including traffic volume (vehicles per interval) and average speed (km/h),  
521 were computed and aggregated into 5-minute intervals. This dataset comprises approximately 1.3  
522 billion data points, providing a comprehensive representation of urban traffic dynamics much larger  
523 than many commonly used open-source benchmarks [28].

524 **Venue-Centric Traffic Data:** To specifically capture traffic patterns influenced by large-scale  
525 public events, supplementary datasets focused on major venues were compiled.

- 526 • *Aggregated Venue Flow:* Derived from location-based service data, aggregated traffic volume  
527 information was obtained for the precise geographical boundaries and surrounding buffer areas  
528 of over 300 major venues (sports stadiums, concert halls, major tourist attractions, transport  
529 hubs) across multiple cities in China. This dataset helps model event-specific demand surges  
530 and dispersion patterns.

- 531     • *Fine-Grained Venue Grid Data*: For a subset of the venues above, traffic volume was aggregated  
532       onto a finer  $100m \times 100m$  grid covering the venue. Due to the high granularity leading to  
533       sparsity, we identified and utilized data primarily from the top-20 grid cells exhibiting the  
534       highest average historical traffic volume within each venue’s defined area, focusing analysis on  
535       the most relevant micro-locations.

536     **Contextual Event Data:** Real-time and historical event information, crucial for the LLM  
537       reasoning stage (Stage 2) and evaluation during anomalies, was primarily sourced from the Amap  
538       open platform APIs and associated historical logs for the relevant periods and geographical areas.  
539       This included:

- 540     • *Structured Construction Data*: A curated dataset encompassing over 10,000 construction events  
541       on highways and major arterials. Each entry typically includes precise location information  
542       (coordinates or road segment identifiers), scheduled start and end times, number and type of  
543       lanes affected (e.g., closure, partial blockage), and nature of the work.
- 544     • *Unstructured Anomaly Reports*: Traffic incident information disseminated via Amap, origi-  
545       nating from user reports or official traffic authority alerts. These reports typically contain a  
546       natural language description of the incident (e.g., "Accident involving two cars on Ring Road  
547       eastbound near Exit 5, blocking right lane"), an approximate location , and a timestamp. This  
548       text serves as direct input for the LLM.

549     **Preprocessing:** Standard preprocessing steps were applied to the traffic datasets before model  
550       training and evaluation. Missing values in the time-series data less than 5% of points were imputed  
551       using linear interpolation while others are dropped. Traffic state features (speed, volume) were  
552       normalized using Z-score normalization based on the mean and standard deviation calculated from  
553       the training portion of the primary Beijing dataset. For GNN-based baselines, the spatial graph  
554       adjacency matrix was constructed based on road network distance threshold, with edge weights  
555       typically defined by inverse distance. Unstructured text data from anomaly reports and event  
556       information was cleaned to remove irrelevant artifacts before being fed into the LLM prompts.

## 557     5.2 Spatio-Temporal Foundation Model for Probabilistic Forecasting

558     The foundation of our Chat-ITS framework is a powerful probabilistic forecaster built upon the  
559       principles of the Chronos framework [21]. This approach reframes forecasting as a language modeling  
560       task. The core idea is to translate a continuous numerical time series into a sequence of discrete  
561       tokens, analogous to words in a sentence, and then use a standard language model architecture to  
562       learn the grammar of traffic dynamics and predict future tokens. For this purpose, we utilized a  
563       T5-based encoder-decoder architecture [22], which we pre-trained on extensive historical traffic data.

564     The process involves three key stages: input tokenization, model training, and probabilistic  
565       forecast generation.

- 566     1. **Time-Series Tokenization:** To make numerical data processable by a language model, each  
567       time series undergoes a two-step tokenization pipeline:

- 568 • **Scaling:** To handle the varying scales of traffic data, each time series is normalized. We  
 569 employ mean scaling, where each value in the series ( $x_t$ ) is divided by the mean of the  
 570 absolute values of its historical context ( $s$ ). This brings diverse series to a comparable  
 571 scale.

$$x'_t = \frac{x_t}{s} \quad (3)$$

- 572 • **Quantization:** The scaled, continuous values ( $x'_t$ ) are then mapped into a finite vocabu-  
 573 lary of  $B$  discrete integer tokens. This is achieved by dividing the range of possible scaled  
 574 values into  $B$  predefined bins. This critical step completes the transformation from a  
 575 sequence of numbers to a sequence of tokens.

576 2. **Model Architecture and Training:** We employ a standard T5 encoder-decoder architec-  
 577 ture, which is based on the Transformer model. The encoder processes the sequence of tokens  
 578 representing historical traffic data to create a rich contextual representation. The core of the  
 579 Transformer is the **Scaled Dot-Product Attention** mechanism, which allows the model to  
 580 weigh the importance of different tokens in the input sequence when making a prediction. The  
 581 attention output is computed as:

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}}\right)\mathbf{V} \quad (4)$$

582 where  $\mathbf{Q}$  (Query),  $\mathbf{K}$  (Key), and  $\mathbf{V}$  (Value) are matrices derived from the input token embed-  
 583 dings, and  $d_k$  is the dimension of the keys. The model further enhances this with Multi-Head  
 584 Attention, running the attention mechanism in parallel multiple times to jointly attend to  
 585 information from different representation subspaces.

586 The decoder then uses the encoder’s output to autoregressively generate the forecast token by  
 587 token. The model is trained to predict the next token by minimizing the negative log-likelihood  
 588 (cross-entropy loss) over the vocabulary of  $B$  bins. The objective is formalized as:

$$L(\theta) = - \sum_{t=T+1}^{T+n} \log P(z_t | z_{<t}; \theta) \quad (5)$$

589 where  $z_t$  is the ground-truth token at a future time step  $t$ ,  $z_{<t}$  represents all preceding tokens,  
 590 and  $P(\cdot)$  is the probability distribution predicted by the model with parameters  $\theta$ . This process  
 591 effectively trains the model to perform “regression via classification.”

592 3. **Probabilistic Forecast Generation:** At inference time, the trained model generates prob-  
 593 abilistic forecasts. For each future time step, the decoder outputs a probability distribution  
 594 across all  $B$  tokens. To capture uncertainty, we generate multiple future scenarios by autore-  
 595 gressively **sampling** from this distribution  $K$  times (in our work,  $K = 20$ ). These sampled  
 596 token sequences are then converted back into numerical trajectories through a reverse pipeline:

- 597 • **De-quantization:** Each token is mapped back to the numerical center of its correspond-  
 598 ing bin.

- 599        • **Re-scaling:** The values are multiplied by the original scaling factor ( $s$ ) to restore them  
600              to their physical scale.

601        This procedure yields a set of  $K$  distinct future trajectories, which collectively form a proba-  
602              bilitistic forecast. This serves as the crucial input for the LLM-based contextual adjustment in  
603              Stage 2 of our framework.

### 604        5.3 LLM Integration Details

605        We primarily utilized *Qwen2.5-14B-Instruct* [23] for the LLM components in Stages 2 and 3, se-  
606              lected for its strong reasoning and instruction-following capabilities. Structured data (weather, road  
607              work) was formatted as key-value pairs. Unstructured text (incidents, events) was included directly,  
608              prefixed with source and time. Time-series candidate trajectories were summarized by providing key  
609              statistics (e.g., min/max/avg speed in critical zones, predicted congestion duration). In the primary  
610              mode, the LLM selected the single most plausible trajectory ( $k^*$ ) from the  $K$  candidates based on  
611              its contextual evaluation:  $\hat{\mathbf{X}}_{T+1:T+n}^* = \hat{\mathbf{X}}_{T+1:T+n}^{(k^*)}$ . In experiments exploring active adjustment, the  
612              LLM’s quantitative impact assessment (e.g., predicted % speed reduction) was used to mathemati-  
613              cally modify the selected or mean trajectory. Reporting/Recommendation Generation: Standard  
614              LLM text generation was used based on the prompts described above. Few-shot examples from the  
615              historical dispatch data were included in the recommendation prompts to bias the LLM towards  
616              operationally relevant suggestions. No LLM fine-tuning was performed for this study. Besides, API  
617              calls were made using standard libraries with default temperature settings (e.g., temperature=0.7)  
618              for generative tasks in Stage 3 to allow for some variability, and lower temperature (e.g., 0.1) for  
619              the selection task in Stage 2 to ensure consistent choices.

### 620        5.4 Baseline Implementations

621        The chosen baselines represent a broad spectrum of modern time series forecasting methodologies,  
622              ensuring a robust and comprehensive comparison against different architectures:

- 623        • DLinear[13]: Represents simple yet surprisingly effective linear models, serving as a strong  
624              benchmark against more complex architectures by decomposing the time series and applying  
625              separate linear layers. It challenges the necessity of intricate designs for certain forecasting  
626              tasks.
- 627        • FiLM [24]: Represents linear models enhanced with frequency analysis, designed to improve  
628              forecasting by better capturing periodicity through specific decomposition techniques applied  
629              in the frequency domain.
- 630        • Informer [25]: A prominent Transformer-based model optimized for long sequence time-series  
631              forecasting (LSTF) efficiency through a ProbSparse self-attention mechanism and distilling  
632              operation, representing efficient Transformer variants.

633 • PatchTST [26]: Represents channel-independent Transformer approaches utilizing patching,  
634 where input time series are divided into subseries-level patches that are fed as tokens to the  
635 Transformer, capturing local semantic information.

636 • Chronos [21]: Represents recent large pre-trained foundation models for time series, leveraging  
637 language model architectures scaled to time series data for zero-shot or few-shot forecasting,  
638 showcasing the potential of large-scale pre-training.

639 • iTransformer [27]: An innovative Transformer architecture that inverts the standard process  
640 by applying attention to embedded variates across the entire time series length, designed to  
641 better capture multivariate correlations.

642 For implementation, we utilized established frameworks such as TSLib [29, 30] or the official  
643 public code repositories associated with each baseline model. To guarantee a fair and direct com-  
644 parison, all baseline models were trained using the identical historical dataset that was employed  
645 for training the Chat-ITS foundation model. The sole exception was Chronos, for which we lever-  
646 aged the publicly available pre-trained weights, applying it directly as a zero-shot forecaster without  
647 fine-tuning on our specific dataset. This curated selection of baselines ensures our evaluation covers  
648 a diverse range of contemporary forecasting architectures, encompassing simple linear approaches,  
649 frequency-domain enhanced models, various Transformer adaptations (including those optimized for  
650 efficiency, employing patching mechanisms, or utilizing inverted attention across variates), and large  
651 pre-trained foundation models.

## 652 5.5 Evaluation Details

653 We evaluated the model’s performance using standard time-series forecasting metrics: Mean Abso-  
654 lute Error (MAE), Mean Squared Error (MSE), Root Mean Squared Error (RMSE), Mean Absolute  
655 Percentage Error (MAPE), and Weighted Absolute Percentage Error (WAPE). MAE measures the  
656 average absolute difference between predictions and actual values. MSE computes the average of  
657 the squared errors, which emphasizes larger deviations more heavily. RMSE is the square root of  
658 MSE, bringing the error back to the original scale of the data. MAPE and WAPE assess the relative  
659 size of errors compared to actual values, providing scale-independent evaluation.

660 It is worth noting that we preferred WAPE over MAPE for traffic volume forecasting tasks, where  
661 the data often contains zero or near-zero values. In such scenarios, MAPE can become undefined or  
662 unstable due to division by zero or extremely small actual values, leading to misleading evaluations.  
663 Additionally, in cases with large variance in traffic volumes, MAPE tends to overemphasize errors in  
664 low-volume periods while underrepresenting high-volume ones. In contrast, WAPE normalizes total  
665 absolute error by the sum of actual values across all time steps and locations, offering a more stable  
666 and representative metric under these conditions.

667 The specific metrics are defined as (6), (7), (8), (9), and (10):

$$\text{MAE} = \frac{1}{nN} \sum_{t=T+1}^{T+n} \sum_{i=1}^N |x_{t,i} - \hat{x}_{t,i}| \quad (6)$$

$$668 \quad \text{MSE} = \frac{1}{nN} \sum t = T + 1^{T+n} \sum_{i=1}^N (x_{t,i} - \hat{x}_{t,i})^2 \quad (7)$$

$$669 \quad \text{RMSE} = \sqrt{\text{MSE}} = \sqrt{\frac{1}{nN} \sum t = T + 1^{T+n} \sum_{i=1}^N (x_{t,i} - \hat{x}_{t,i})^2} \quad (8)$$

$$670 \quad \text{MAPE} = \frac{1}{nN} \sum t = T + 1^{T+n} \sum_{i=1}^N \left| \frac{x_{t,i} - \hat{x}_{t,i}}{x_{t,i}} \right| \quad (9)$$

$$671 \quad \text{WAPE} = \frac{\sum_{t=T+1}^{T+n} \sum_{i=1}^N |x_{t,i} - \hat{x}_{t,i}|}{\sum t = T + 1^{T+n} \sum_{i=1}^N |x_{t,i}|} \quad (10)$$

672     **Anomaly Identification:** For evaluating performance during anomalies, event periods were  
 673     identified using timestamps from the Amap/Gaode construction and incident logs. Anomalous  
 674     periods were defined as 1 hour before to 1 hours after the logged event time for relevant locations.  
 675     For public events, the anomalous period covered 2 hours before the event start to 2 hours after  
 676     the event end. Zero-shot evaluation used events from categories completely held out during any  
 677     training/fine-tuning.

678     **Data Splits:** Data was split chronologically for each city/dataset. Typically, the first 70% was  
 679     used for pre-training the foundation model, the next 10% for validation (e.g., selecting foundation  
 680     model checkpoints, basic prompt tuning), and the final 20% for testing.

## 680     **Acknowledgments**

681     This work was supported by Beijing Natural Science Foundation (No. JQ24051), Beijing Nova  
 682     Program (No. 20230484432) and Independent Research Project of the State Key Laboratory of  
 683     Intelligent Green Vehicle and Mobility, Tsinghua University (No. ZZ-GG-20250406).

## 684     **6 Data Availability**

685     The datasets central to this study, including the large-scale traffic network data and the anomalous  
 686     event logs, were provided by Amap. Due to the proprietary nature of this information, which encom-  
 687     passes commercial sensitivities and privacy considerations, these datasets are not publicly available.  
 688     We acknowledge the importance of reproducibility and regret that these necessary restrictions pre-  
 689     vent the public dissemination of these materials. Enquiries regarding the methodology or potential  
 690     collaborations may be directed to the corresponding author.

## 691     **7 Supplementary Materials**

- 692         • **Text S1 and Table S1.** Detailed description and illustrative examples of the Structured  
 693         Construction Data used in this study.

- 694 • **Text S2 and Table S2.** Detailed description and examples of the Unstructured Anomaly  
695 Reports.

696 **References**

- 697 1. Wu K, Ding J, Lin J, et al. Big-data empowered traffic signal control could reduce urban carbon  
698 emission. *Nature Communications* 2025;16. Publisher: Nature Publishing Group:2013.
- 699 2. Avila AM and Mezić I. Data-driven analysis and forecasting of highway traffic dynamics. *Nature*  
700 *Communications* 2020;11. Publisher: Nature Publishing Group TLDR: Here it is demonstrated  
701 how the Koopman mode decomposition can offer a model-free, data-driven approach for ana-  
702 lyzing and forecasting traffic dynamics.:2090.
- 703 3. Shao Z, Zhang Z, Wang F, and Xu Y. Pre-training Enhanced Spatial-temporal Graph Neural  
704 Network for Multivariate Time Series Forecasting. In: *Proceedings of the 28th ACM SIGKDD*  
705 *Conference on Knowledge Discovery and Data Mining*. 2022:1567–77. doi: 10.1145/3534678.  
706 3539396. arXiv: 2206.09113[cs]. URL: <http://arxiv.org/abs/2206.09113> (visited on  
707 10/24/2023).
- 708 4. Li Z, Cai R, Fu TZJ, Hao Z, and Zhang K. Transferable Time-Series Forecasting Under Causal  
709 Conditional Shift. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 2024;46.  
710 Conference Name: IEEE Transactions on Pattern Analysis and Machine Intelligence:1932–49.
- 711 5. Runge J, Bathiany S, Boltt E, et al. Inferring causation from time series in Earth system  
712 sciences. *Nature Communications* 2019;10. Publisher: Nature Publishing Group:2553.
- 713 6. Yuan Y, Ding J, Feng J, Jin D, and Li Y. UniST: A Prompt-Empowered Universal Model  
714 for Urban Spatio-Temporal Prediction. 2024. doi: 10.1145/3637528.3671662. arXiv: 2402.  
715 11838[cs]. URL: <http://arxiv.org/abs/2402.11838> (visited on 07/15/2024).
- 716 7. Guo X, Zhang Q, Jiang J, Peng M, Zhu M, and Yang HF. Towards explainable traffic flow  
717 prediction with large language models. *Communications in Transportation Research* 2024;4.  
718 TLDR: This paper proposes a Traffic flow Prediction model based on Large Language Models  
719 (LLMs) to generate explainable traffic predictions, named xTP-LLM, and is the first study to  
720 use LLM for explainable prediction of traffic flows.:100150.
- 721 8. Williams AR, Ashok A, Marcotte É, et al. Context is Key: A Benchmark for Forecasting  
722 with Essential Textual Information. TLDR: A time series forecasting benchmark that pairs  
723 numerical data with diverse types of carefully crafted textual context, requiring models to  
724 integrate both modalities, and a simple yet effective LLM prompting method that outperforms  
725 all other tested methods on this benchmark. 2024. doi: 10.48550/arXiv.2410.18959. arXiv:  
726 2410.18959[cs]. URL: <http://arxiv.org/abs/2410.18959> (visited on 01/08/2025).

- 727 9. Liu H, Xu S, Zhao Z, et al. Time-MMD: A New Multi-Domain Multimodal Dataset for Time  
728 Series Analysis. TLDR: Time-MMD is introduced, the first multi-domain, multimodal time  
729 series dataset covering 9 primary data domains, and MM-TSFlib, the first multimodal time-  
730 series forecasting (TSF) library, seamlessly pipelining multimodal TSF evaluations based on  
731 Time-MMD for in-depth analyses. 2024. DOI: 10.48550/arXiv.2406.08627. arXiv: 2406.  
732 08627[cs]. URL: <http://arxiv.org/abs/2406.08627> (visited on 07/26/2024).
- 733 10. Li Z, Lin X, Liu Z, et al. Language in the Flow of Time: Time-Series-Paired Texts Weaved  
734 into a Unified Temporal Narrative. 2025. DOI: 10.48550/arXiv.2502.08942. arXiv: 2502.  
735 08942[cs]. URL: <http://arxiv.org/abs/2502.08942> (visited on 02/17/2025).
- 736 11. Xue H and Salim FD. PromptCast: A New Prompt-based Learning Paradigm for Time Series  
737 Forecasting. 2023. DOI: 10.48550/arXiv.2210.08964. arXiv: 2210.08964[cs, math, stat].  
738 URL: <http://arxiv.org/abs/2210.08964> (visited on 01/17/2024).
- 739 12. Arango SP, Mercado P, Kapoor S, et al. ChronosX: Adapting Pretrained Time Series Mod-  
740 els with Exogenous Variables. TLDR: This paper introduces a new method to incorporate  
741 covariates into pretrained time series forecasting models through modular blocks that inject  
742 past and future covariate information, without necessarily modifying the pretrained model  
743 in consideration. 2025. DOI: 10.48550/arXiv.2503.12107. arXiv: 2503.12107[cs]. URL:  
744 <http://arxiv.org/abs/2503.12107> (visited on 03/27/2025).
- 745 13. Zeng A, Chen M, Zhang L, and Xu Q. Are Transformers Effective for Time Series Forecast-  
746 ing? 2022. arXiv: 2205.13504[cs]. URL: <http://arxiv.org/abs/2205.13504> (visited on  
747 07/25/2023).
- 748 14. Tan M, Merrill MA, Gupta V, Althoff T, and Hartvigsen T. Are Language Models Actually  
749 Useful for Time Series Forecasting? 2024. DOI: 10.48550/arXiv.2406.16964. arXiv: 2406.  
750 16964[cs]. URL: <http://arxiv.org/abs/2406.16964> (visited on 08/06/2024).
- 751 15. Yuan X and Qiao Y. Diffusion-TS: Interpretable Diffusion for General Time Series Generation.  
752 In: The Twelfth International Conference on Learning Representations. 2023. URL: <https://openreview.net/forum?id=4h1apFj099> (visited on 12/13/2024).
- 754 16. Su J, Jiang C, Jin X, et al. Large Language Models for Forecasting and Anomaly Detection:  
755 A Systematic Literature Review. 2024. DOI: 10.48550/arXiv.2402.10350. arXiv: 2402.  
756 10350[cs]. URL: <http://arxiv.org/abs/2402.10350> (visited on 04/02/2024).
- 757 17. Wang X, Fang M, Zeng Z, and Cheng T. Where Would I Go Next? Large Language Mod-  
758 els as Human Mobility Predictors. 2023. DOI: 10.48550/arXiv.2308.15197. arXiv: 2308.  
759 15197[physics]. URL: <http://arxiv.org/abs/2308.15197> (visited on 10/30/2023).
- 760 18. Feng J, Du Y, Liu T, Guo S, Lin Y, and Li Y. CityGPT: Empowering Urban Spatial Cognition  
761 of Large Language Models. 2024. DOI: 10.48550/arXiv.2406.13948. arXiv: 2406.13948[cs].  
762 URL: <http://arxiv.org/abs/2406.13948> (visited on 08/01/2024).
- 763 19. Liu P, Guo H, Dai T, et al. Taming Pre-trained LLMs for Generalised Time Series Forecasting  
764 via Cross-modal Knowledge Distillation. 2024. DOI: 10.48550/arXiv.2403.07300. arXiv:  
765 2403.07300[cs]. URL: <http://arxiv.org/abs/2403.07300> (visited on 03/19/2024).

- 766 20. Jin M, Wang S, Ma L, et al. Time-LLM: Time Series Forecasting by Reprogramming Large  
767 Language Models. 2023. DOI: 10.48550/arXiv.2310.01728. arXiv: 2310.01728[cs]. URL:  
768 <http://arxiv.org/abs/2310.01728> (visited on 10/19/2023).
- 769 21. Ansari AF, Stella L, Turkmen C, et al. Chronos: Learning the Language of Time Series. 2024.  
770 arXiv: 2403.07815[cs]. URL: <http://arxiv.org/abs/2403.07815> (visited on 03/18/2024).
- 771 22. Raffel C, Shazeer N, Roberts A, et al. Exploring the limits of transfer learning with a unified  
772 text-to-text transformer. Journal of machine learning research 2020;21:1–67.
- 773 23. Bai J, Bai S, Chu Y, et al. Qwen Technical Report. arXiv preprint arXiv:2309.16609 2023.
- 774 24. Zhou T, Ma Z, Wen Q, et al. Film: Frequency improved legendre memory model for long-term  
775 time series forecasting. Advances in neural information processing systems 2022;35:12677–90.
- 776 25. Zhou H, Zhang S, Peng J, et al. Informer: Beyond efficient transformer for long sequence time-  
777 series forecasting. In: *Proceedings of the AAAI conference on artificial intelligence*. Vol. 35. 12.  
778 2021:11106–15.
- 779 26. Nie Y, Nguyen NH, Sinthong P, and Kalagnanam J. A Time Series is Worth 64 Words: Long-  
780 term Forecasting with Transformers. 2023. arXiv: 2211.14730[cs]. URL: <http://arxiv.org/abs/2211.14730> (visited on 12/15/2023).
- 782 27. Liu Y, Hu T, Zhang H, et al. iTransformer: Inverted Transformers Are Effective for Time  
783 Series Forecasting. 2023. DOI: 10.48550/arXiv.2310.06625. arXiv: 2310.06625[cs]. URL:  
784 <http://arxiv.org/abs/2310.06625> (visited on 10/24/2023).
- 785 28. Liu X, Xia Y, Liang Y, et al. LargeST: A Benchmark Dataset for Large-Scale Traffic Fore-  
786 casting. 2023. arXiv: 2306.08259[cs]. URL: <http://arxiv.org/abs/2306.08259> (visited on  
787 10/31/2023).
- 788 29. Wu H, Hu T, Liu Y, Zhou H, Wang J, and Long M. TimesNet: Temporal 2D-Variation Modeling  
789 for General Time Series Analysis. In: *International Conference on Learning Representations*.  
790 2023.
- 791 30. Wang Y, Wu H, Dong J, Liu Y, Long M, and Wang J. Deep Time Series Models: A Compre-  
792 hensive Survey and Benchmark. 2024.